

Linux 铁三角之I/O(一)

Yomocode.com

麦当劳喜欢您来，喜欢您再来



扫描关注
Linux阅码场




I/O与网络模型

- *阻塞
- *非阻塞
- *多路复用
- *Signal IO
- *异步IO
- *libevent

yomocode 阅码场


场景案例

- 初始化，而后开背景图片



场景案例


- CPU算包，网卡发包



Bootchart


Boot chart for Android (01/01/00 08:06:04)

uname: Linux version 3.0.36+ (tjin@VLC-INTEL) (gcc version 4.6.x-google 20120106 (prerelease) (GCC)) #1 SMP PREEMPT Mon Apr 28 10:10:
release: 0.0
CPU: ARMv7 Processor rev 0 (v7l)
kernel options: console=ttyFIQ0 androidboot.console=ttyFIQ0 init=/init initrd=0x62000000,0x00130000 mtdparts=rk29xxnand:0x00002000@0x00002000(misc),0x00004000@0x00004000(kernel),0x00008000@0x00008000(boot),0x00010000@0x00010000(recovery)
time: 1:41



Systemd readahead

- `systemd-readahead-collect.service` 搜集系统启动过程中的文件访问信息; `systemd-readahead-replay.service` 在后续启动过程中完成回放。




把CPU和I/O的交替等，变为CPU和I/O的同时工作，充分利用系统资源。


阻塞

进程阻塞等I/O ready

read你开始()



read返回



触摸屏没按?

触摸屏按了!


EINTR

被信号打断的系统调用 Interrupted system call

```
act.sa_handler = sig_handler;
act.sa_flags = 0;
// act.sa_flags |= SA_RESTART;
sigemptyset(&act.sa_mask);
if (-1 == sigaction(SIGUSR1, &act, &oldact)) {
...
do {
    ret = read(STDIN_FILENO, buf, 10);
    if ((ret == -1) && (errno == EINTR))
        printf("retry after eintr\n");
} while(((ret != -1) && (errno == EINTR));)
```

非阻塞


进程不会等I/O ready



多路复用

```
int select(int nfds, fd_set *readfds, fd_set *writefds,  
           fd_set *exceptfds, struct timeval *timeout);
```

select()开始



进程睡眠等待

所有fd都不满足
读写条件

select()返回 read开始()



有一个fd满足
读写条件

epoll

Epoll的事件注册

```
int epoll_ctl(int epfd, intop, int fd, struct epoll_event*event);
```

EPOLL_CTL_ADD: 注册新的fd到epfd中;

EPOLL_CTL_MOD: 修改已经注册的fd的
监听事件;

EPOLL_CTL_DEL: 从epfd中删除一个fd;

等待事件触发

```
int epoll_wait(int epfd, struct epoll_event * events,  
intmaxevents, int timeout);
```

SIGIO app

执行signal handler



SIGIO

触摸屏按了!

内核

异步I/O

Glibc-AIO

`aio_read`

立即返回

中间干别的事情

`aio_suspend`

后台线程读

以同步原语等待读结束

Kernel-AIO w/o _DIRECT

准备上下文

```
int io_setup(unsigned nr_events,  
             aio_context_t *ctx_idp);
```

发布io请求

```
int io_submit(aio_context_t ctx_id, long nr,  
             struct iocb **iocbpp);
```

等待Completions

```
int io_getevents(aio_context_t ctx_id, long  
                min_nr, long nr, struct io_event *events,  
                struct timespec *timeout);
```

销毁上下文

```
int io_destroy(aio_context_t ctx_id);
```

Libevent

libevent是一个Reactor:

一种事件驱动机制，注册回调函数，如果事件发生，被回调。

```
static void cmd_event(int fd, short events, void *arg)
{
    char msg[1024];

    int ret = read(fd, msg, sizeof(msg));
    if (ret <= 0) {
        perror("read fail");
        exit(1);
    }
    msg[ret]='\0';
    printf("%s", msg);
}

int main(int argc, char **argv)
{
    struct event ev_cmd;
    event_init();
    //监听终端输入事件
    event_set(&ev_cmd, STDIN_FILENO,
              EV_READ | EV_PERSIST, cmd_event, NULL);

    event_add(&ev_cmd, NULL);
    event_dispatch();

    return 0;
}
```

有人输入，触发回调

C10K 问题

大多数开发人员都能很容易地从功能上实现，
但一旦放到大并发场景下.....

← → ↻ ⓘ www.kegel.com/c10k.html

The C10K problem

[\[Help save the best Linux news source on the web - subscribe to Linux Weekly News!\]](#)

It's time for web servers to handle ten thousand clients simultaneously, don't you think? After all, the we

And computers are big, too. You can buy a 1000MHz machine with 2 gigabytes of RAM and an 1000Mb clients, that's 50KHz, 100Kbytes, and 50Kbits/sec per client. It shouldn't take any more horsepower than the network once a second for each of twenty thousand clients. (That works out to \$0.08 per client, by tl systems charge are starting to look a little heavy!) So hardware is no longer the bottleneck.

In 1999 one of the busiest ftp sites, cdrom.com, actually handled 10000 clients simultaneously through a [being offered by several ISPs](#), who expect it to become increasingly popular with large business custom

And the thin client model of computing appears to be coming back in style -- this time with the server c

With that in mind, here are a few notes on how to configure operating systems and write code to supp Unix-like operating systems, as that's my personal area of interest, but Windows is also covered a bit.

Contents

模型对比

模型	特点
一个连接，一个进程/线程	进程/线程会占用大量的系统资源，切换开销大；可扩展性差
一个进程/线程，处理多个连接 select	fd上限+重复初始化+逐个排查所有fd状态， $O(n)$ 的效率不断地去查fd
一个进程/线程，处理多个连接 epoll	epoll_wait()返回的时候只给应用提供发生了状态变化的fd 典型用户：nginx, node.js
Libevent: 跨平台，封装底层平台的调用，提供统一的 API (Windows-IOCP, Solaris- /dev/poll, FreeBSD-kqueue, Linux - epoll)	当一个fd的特定事件（如可读，可写或出错）发生了，libevent就会自动执行用户指定的callback，来处理事件。

谢谢！

yomocode 阅码场