

Linux 铁三角之I/O(三)

麦当劳喜欢您来，喜欢您再来



扫描关注
Linux阅码场



文件系统的实现

- *EXT2/3/4的layout
- *文件系统的一致性:append一个文件的全流程
- *掉电与文件系统一致性
- *fsck
- *文件系统的日志
- *ext4 mount选项
- *文件系统的debug和dump
- *Copy On Write文件系统: btrfs


- Super block, storing:
 - Size and location of bitmaps
 - Number and location of inodes
 - Number and location of data blocks
 - Index of root inodes

Bitmap of free & used data blocks

Bitmap of free & used inodes\

- Table of inodes
- Each inode is a file/directory
- Includes meta-data and lists of associated data blocks

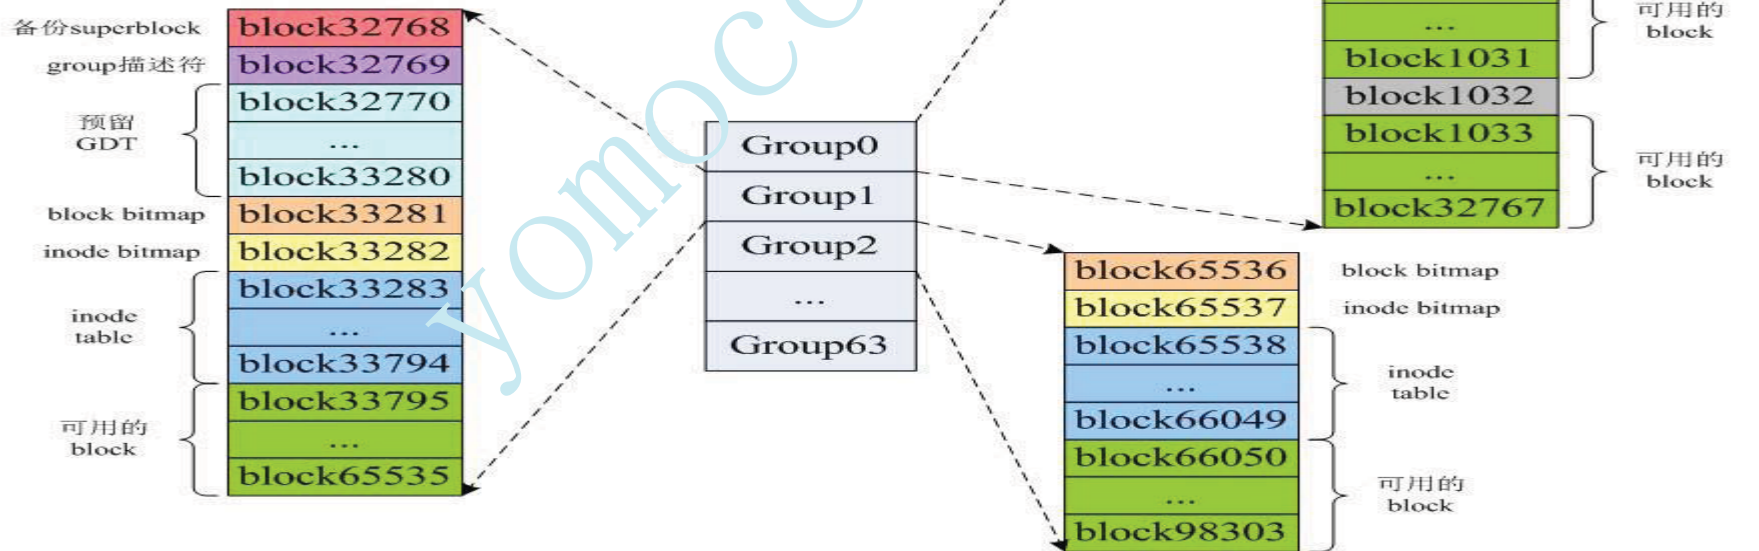
Data blocks (4KB each)




Group

superblock 记录文件系统的类型、block大小、block总数、free block数、inode大小、inode总数、free inode数、group的总数等,在多个group进行备份。


Group描述符记录: block bitmap位置、inode bitmap位置、free block, free inode数量。




Append file 一个文件



文件系统必须完成三件事：inode ($I[v2]$), bitmap ($B[v2]$)和data block (Db).
crash 也许中途发生，并没有完成三件事，文件系统在funny的状态。



硬件不可能原子执行



掉电与文件系统一致性

1. 任何的软件技术都无法保证掉电不丢数据！
只能保证一致性(元数据+数据的一致性，或者仅元数据的一致性)！
2. `dirty_expire_centisecs`、`DIRECT_IO`、`SYNC IO`的调整，不影响丢或者不丢数据，只影响丢多少数据。
3. `fsck`、日志、`CoW`文件系统等技术，帮忙提供一致性。


fsck (file system consistency check)

unclean shutdown后自动运行，或者手动运行；检查**superblock**、**inode**和**free block bitmap**、所有**inode**的**Reachability**（比如删除**corrupted**的**inode**）、验证目录的一致性；

Dennis Ritchie: "So fsck was originally called something else"


Question: "What was it called?"

Dennis Ritchie: "Well, the **second** letter was different"



日志：transaction


重启的过程中：transactions that are committed but not free are replayed in order



- ✓ The transaction begin (TxB) tells us about this update, including information about the pending update to the file system (e.g., the final addresses of the blocks I[v2], B[v2], and Db), as well as some kind of transaction identifier (TID).
- ✓ The middle three blocks just contain the exact contents of the blocks themselves;
- ✓ The final block (TxE) is a marker of the end of this transaction, and will also contain the TID.

Commits 和 Checkpoints


- 日志的写完成, transaction committed
- Transaction完成, OS checkpoint这次更新




- Final step: free the checkpointed transaction

日志的4个阶段


- ✓1. Journal write: Write the contents of the transaction (including TxB, metadata, and data) to the log; wait for these writes to complete.



- ✓2. Journal commit: Write the transaction commit block (containing TxE) to the log; wait for write to complete; transaction is said to be committed.



- ✓3. Checkpoint: Write the contents of the update (metadata and data) to their final on-disk locations.



- ✓4. Free: Some time later, mark the transaction free in the journal by updating the journal superblock.

Metadata 日志 - 5 个阶段

完整的日志(`data=journal`)让系统变得很慢

- ✓ 1. Data write: Write data to final location; wait for completion (the wait is optional: `data=writeback/ data=ordered`).
- ✓ 2. Journal metadata write: Write the begin block and metadata to the log; wait for writes to complete.
- ✓ 3. Journal commit: Write the transaction commit block (containing TxE) to the log; wait for the write to complete; the transaction (including data) is now committed.



- ✓ 4. Checkpoint metadata: Write the contents of the metadata update to their final locations within the file system.
- ✓ 5. Free: Later, mark the transaction free in journal superblock.

常用工具

mkfs

dumpe2fs

blkcat

dd

debugfs:

 icheck: block号找inode

 ncheck: inode号找文件名

blktrace


```
baohua@baohua-VirtualBox:~/develop/linux$ sudo blktrace -d /dev/sda -o - | blkparse -i -  
[sudo] password for baohua:
```

设备号 Major:minor	CPU	序号	时间戳	PID	事件	开始块+块数	进程名
8,0	0	1	0.000000000	167	A WS	12957520 + 8 <- (8,1) 12955472	
8,0	0	2	0.000002319	167	Q WS	12957520 + 8 [jbd2/sda1-8]	
8,0	0	3	0.000007055	167	G WS	12957520 + 8 [jbd2/sda1-8]	
8,0	0	4	0.000008384	167	P N	[jbd2/sda1-8]	
8,0	0	5	0.000010005	167	A WS	12957528 + 8 <- (8,1) 12955480	
8,0	0	6	0.000010513	167	Q WS	12957528 + 8 [jbd2/sda1-8]	
8,0	0	7	0.000012040	167	M WS	12957528 + 8 [jbd2/sda1-8]	
8,0	0	8	0.000013173	167	A WS	12957536 + 8 <- (8,1) 12955488	
8,0	0	9	0.000013638	167	Q WS	12957536 + 8 [jbd2/sda1-8]	
8,0	0	10	0.000014240	167	M WS	12957536 + 8 [jbd2/sda1-8]	
8,0	0	11	0.000015133	167	A WS	12957544 + 8 <- (8,1) 12955496	
8,0	0	12	0.000015596	167	Q WS	12957544 + 8 [jbd2/sda1-8]	
8,0	0	13	0.000016182	167	M WS	12957544 + 8 [jbd2/sda1-8]	
8,0	0	14	0.000016963	167	A WS	12957552 + 8 <- (8,1) 12955504	
8,0	0	15	0.000017427	167	Q WS	12957552 + 8 [jbd2/sda1-8]	
8,0	0	16	0.000018613	167	M WS	12957552 + 8 [jbd2/sda1-8]	
8,0	0	17	0.000019205	167	A WS	12957560 + 8 <- (8,1) 12955512	
8,0	0	18	0.000019772	167	Q WS	12957560 + 8 [jbd2/sda1-8]	
8,0	0	19	0.000020361	167	M WS	12957560 + 8 [jbd2/sda1-8]	
8,0	0	20	0.000021561	167	A WS	12957568 + 8 <- (8,1) 12955520	
8,0	0	21	0.000022025	167	Q WS	12957568 + 8 [jbd2/sda1-8]	
8,0	0	22	0.000022608	167	M WS	12957568 + 8 [jbd2/sda1-8]	
8,0	0	23	0.000023381	167	A WS	12957576 + 8 <- (8,1) 12955528	

读写, sync, barrier

COW文件系统:btrfs

每次写磁盘时，先将更新数据写入一个新的 block，当新数据写入成功之后，再更新相关的数据结构指向新 block。
没有日志，用COW实现文件系统一致性。



Transactions in Btrfs

- There is no journal as extN, or xfs has
- COW is used to guarantee consistency
- On fs tree or extent tree modification
 - 1 Tree is cloned and the branch in question is copied and modified
 - 2 New roots are added into root tree (identified by transaction ID)
 - 3 New root of root tree added into superblock
 - 4 Wait on all respective data and metadata to hit the disk
 - 5 Commit the superblock to the disk
 - 6 Original trees can be freed (decrease refcount)
- In case of crash, the original root tree is used (the one in the most up-to-date superblock)

Btrfs: subvolume 和 snapshot

subvolume 可以单独挂载

```
#btrfs subvolume create sub1  
#btrfs subvolume list .  
#mount -o subvolid=256 image aaa
```

snapshot 类似 git branch，也具备 subvolume 特征

```
#btrfs subvolume snapshot snapshot1  
#btrfs subvolume list .  
#mount -o subvolid=257 image aaa  
#btrfs subvolume set-default 257 snapshot1
```

谢谢！

yomocode 阅码场