

Linux任督二脉之内存管理(五)

讲解时间：6月10-14日晚9点
宋宝华 <21cnbao@gmail.com>

微信群直播

扫描二维码报名



Linux任督二脉

(学习形式：微信群)



© 2015 Linux

麦当劳喜欢您来，喜欢您再来



扫描关注
Linux阅码场



其他工程问题以及调优

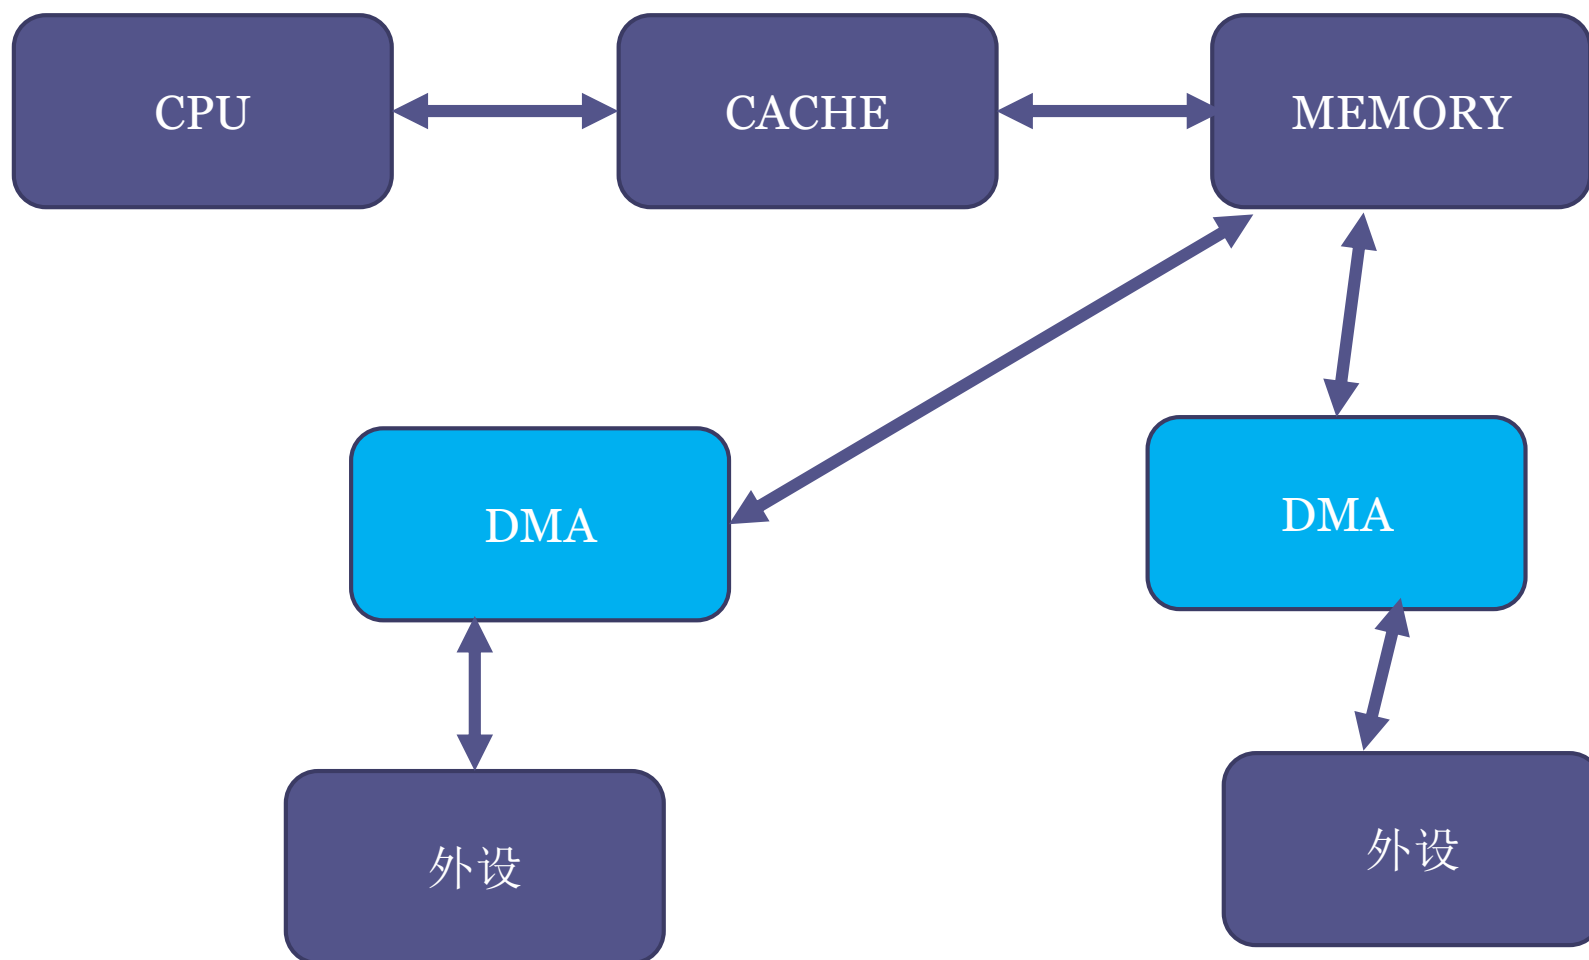
- *DMA和cache一致性
- *内存的cgroup
- *性能方面的调优: page in/out, swapin/out
- *Dirty ratio的一些设置
- *swappiness

练习题

- *vmstat;
- *smem -s swap -t -k -n
- *限制一个cgroup的memory, 用cgexec把一个进程放到这个cgroup

DMA与Cache一致性

- DMA传输外设数据到memory，cache中可能是老的数据
- CPU写数据到memory，cache中是新数据，MEM是老数据



DMA APIs

■ Coherent DMA buffers

```
void * dma_alloc_coherent(struct device *dev, size_t size, dma_addr_t *dma_handle,  
gfp_t flag);
```

```
void dma_free_coherent (struct device *dev, size_t size,  
void *cpu_addr, dma_addr_t dma_handle);
```

CMA和此API自动融合，调用dma_alloc_coherent()将从CMA获得内存。

■ DMA Streaming Mapping

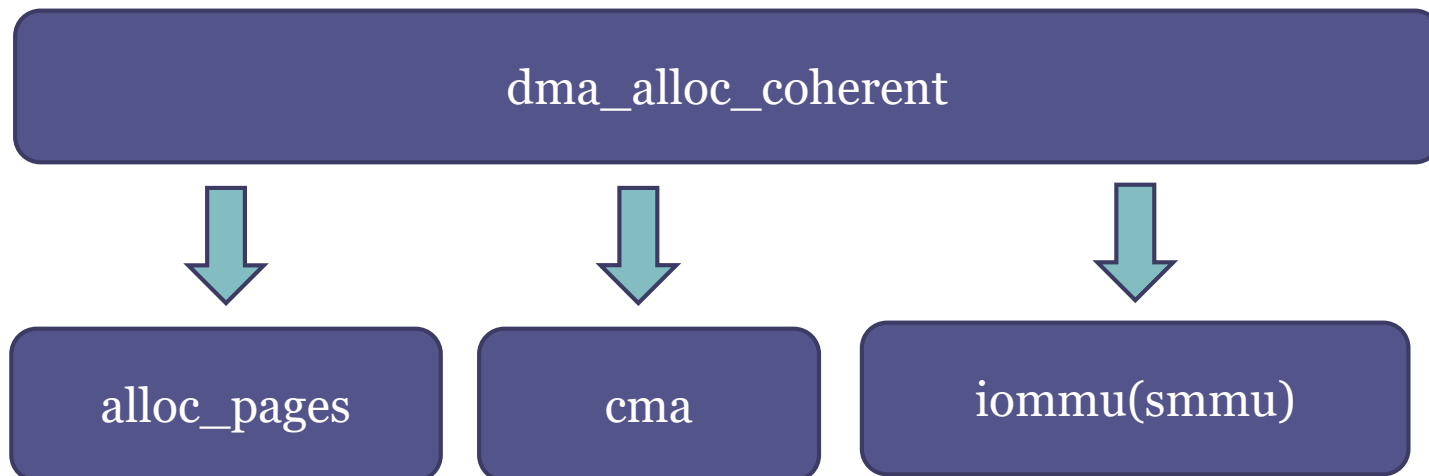

```
int dma_map_sg(...);  
void dma_unmap_sg (...);  
dma_addr_t dma_map_single(...);  
void dma_unmap_single (...);
```

iommu(smmu)

- 把不连续的内存用来做DMA,以及限制DMA的访问范围（protection）
- IOMMU被融入到DMA APIs

I see no reason why ARM should be any different from other architectures which have IOMMUs, and I don't see why ARM should have to invent a whole new framework to handle IOMMUs. And I see no explanation why the existing hooks are unsuitable - at least as the `_initial_` starting point.

-Russell King



cgroup

- 在Linux读写文件时，它用于缓存文件的逻辑内容，从而加快对磁盘上映像和数据的访问

设置cgroup A的最大内存为200MB

```
root@baohua-VirtualBox:~# cd /sys/fs/cgroup/memory/  
root@baohua-VirtualBox:/sys/fs/cgroup/memory# mkdir A  
root@baohua-VirtualBox:/sys/fs/cgroup/memory# cd A/  
root@baohua-VirtualBox:/sys/fs/cgroup/memory/A# echo $((200*1024*1024)) > memory.limit_in_bytes
```

把进程放入cgroup A执行

```
root@baohua-VirtualBox:~/develop/training/memory-courses/day2# cgexec -g memory:A ./a.out  
malloc buffer: 0x3a590008  
0MB written  
...  
192MB written  
196MB written  
Killed
```

文件 Dirty 数据的写回

- `dirty_ratio`

the number of pages at which a process which is generating disk writes will itself start writing out dirty data.

- `dirty_expire_centisecs`

This tunable is used to define when dirty data is old enough to be eligible for writeout by the kernel flusher threads. It is expressed in 100'ths of a second. Data which has been dirty in-memory for longer than this interval will be written out next time a flusher thread wakes up.

- `dirty_writeback_centisecs`

The kernel flusher threads will periodically wake up and write 'old' data out to disk. Setting this to zero disables periodic writeback altogether.

- `dirty_background_ratio`

the number of pages at which the background kernel flusher threads will start writing out dirty data.

vfs_cache_pressure

- 该文件表示内核回收用于directory和inode cache内存的倾向

This variable controls the tendency of the kernel to reclaim the memory which is used for caching of VFS caches, versus pagecache and swap. Increasing this value increases the rate at which VFS caches are reclaimed. It is difficult to know when this should be changed, other than by experimentation. The slabtop command (part of the package procs) shows top memory objects used by the kernel. The vfs caches are the "dentry" and the "*_inode_cache" objects. If these are consuming a large amount of memory in relation to pagecache, it may be worth trying to increase pressure. Could also help to reduce swapping. The default value is 100.

水位设置:min_free_kbytes与lowmem_reserve_ratio

■ min_free_kbytes

This is used to force the Linux VM to keep a minimum number of kilobytes free. The VM uses this number to compute a watermark[WMARK_MIN] value for each lowmem zone in the system. Each lowmem zone gets a number of reserved free pages based proportionally on its size.

PF_MEMALLOC: 紧急内存，可以忽略内存管理的水印进行分配

- ✓ **min_free_kbytes** = $4 * \text{sqrt}(\text{lowmem_kbytes})$, **min_free_kbytes**随着内存的增大不是线性增
- ✓ $\text{watermark}[\text{min}] = \text{per_zone_min_free_pages}$
- ✓ $\text{watermark}[\text{high}] - \text{watermark}[\text{low}] = \text{watermark}[\text{low}] - \text{watermark}[\text{min}] = \text{per_zone_min_free_pages} * 1/4$

■ lowmem_reserve_ratio

The 'lowmem_reserve_ratio' tunable determines how aggressive the kernel is in defending these lower zones.

水位设置:high,low,min

— high:

内存到此点，停止回收

— low: 内存到此点，kswapd启动reclaim

Direct reclaim:直接在应用程序的进程上下文中进行回收，
会阻塞应用

— min: 内存到此点，做**direct reclaim**

```
/proc/sys/vm$ cat /proc/zoneinfo
Node 0, zone    DMA
  pages free    1228
    min         197
    low         246
    high        295
```

```
...
Node 0, zone    Normal
  pages free    21367
    min        10842
    low        13552
    high       16263
```

```
...
/proc/sys/vm$ cat min_free_kbytes
44160
```

swappiness

- Swappiness反映是否积极地使用swap空间

✓ swappiness = 0

仅在内存不足的情况下([free and file-backed pages < high water mark in a zone](#)), 使用swap空间

✓ swappiness = 60

默认值

✓ swappiness = 100

内核将积极的使用swap空间。

```
root@baohua-VirtualBox:/proc/sys/vm# cat swappiness  
60
```

getdelays

- Documentation/accounting/getdelays.c工具,测量调度、I/O、SWAP、Reclaim的延迟

```
root@baohua-VirtualBox:~/develop/linux/Documentation/accounting# gcc getdelays.c
root@baohua-VirtualBox:~/develop/linux/Documentation/accounting# ./a.out -dilv -p 1
print delayacct stats ON
printing IO accounting
listen forever
debug on
family id 25
Sent pid/tgid, retval 0
received 364 bytes
nlmsg_hdr size=16, nlmsg_len=364, rep_len=364
PID      1
```

CPU	count	real total	virtual total	delay total	delay average
	3748	1692000000	1777525852	500993577	0.134ms
IO	count	delay total	delay average		
	194	116602567		0ms	
SWAP	count	delay total	delay average		
	2	731594		0ms	
RECLAIM	count	delay total	delay average		
	0	0		0ms	

init: read=6844416, write=0, cancelled_write=0
^C

vmstat

- **vmstat**可以展现给定时间间隔的服务器的状态值,包括Linux的CPU使用率,内存使用,虚拟内存交换情况,IO读写情况。

```
baohua@baohua-VirtualBox:~$ vmstat 1
```

procs		-----memory-----				---swap--		-----io-----		system--		-----cpu-----				
r	b	swpd	free	buff	cache	si	so	bi	bo	in	cs	us	sy	id	wa	st
1	0	0	126444	361244	61208	1	4	43	5	23	101	2	0	98	0	0
1	0	0	126440	361252	61212	0	0	0	16	990	4430	27	3	70	0	0
1	0	0	125680	361252	61232	0	0	0	0	987	4753	28	2	70	0	0
1	0	0	125204	361252	61248	0	0	0	0	1074	4533	28	2	69	0	0
2	0	0	125292	361252	61260	0	0	0	0	998	4380	27	3	71	0	0
1	0	0	125784	361252	61272	0	0	0	0	913	4915	28	2	70	0	0
1	0	0	125844	361260	61276	0	0	0	332	995	4118	28	2	70	0	0
2	0	0	124996	361260	61296	0	0	0	0	972	4415	26	3	71	0	0
1	0	0	124992	361260	61308	0	0	0	0	978	5045	28	3	70	0	0
2	0	0	125060	361260	61320	0	0	0	0	964	4229	27	2	71	0	0
1	0	0	125072	361260	61332	0	0	0	0	993	5155	29	2	69	0	0

课程练习源码

<https://github.com/21cnbao/memory-courses>

课后阅读

https://access.redhat.com/documentation/en-us/red_hat_enterprise_linux/6/html/performance_tuning_guide/s-memory-tunables

更早课程

- 《Linux总线、设备、驱动模型》录播：

<http://edu.csdn.net/course/detail/5329>

- 深入探究Linux的设备树

<http://edu.csdn.net/course/detail/5627>

- C语言大型软件设计的面向对象

<https://edu.csdn.net/course/detail/6496>

谢谢！