

Linux 任务二 内存管理(二)

Linux 阅码场 yomocode.com

麦当劳喜欢您来，喜欢您再来



扫描关注
Linux阅码场



内存的动态申请和释放

- * slab、kmalloc/kfree、/proc/slabinfo和slabtop
- * 用户空间malloc/free与内核之间的关系
- * mallopt
- * vmalloc
- * 内存耗尽（OOM）、oom_score和oom_adj
- * Android进程生命周期与OOM

练习题

- *看/proc/slabinfo，运行slabtop
- *运行mallopt.c程序：mallopt(M_TRIM_THRESHOLD)等
- *看/proc/vmallocinfo,grep ioremap映射
- *运行一个很耗费内存的程序，观察oom memory
- *通过oom_adj调整firefox的oom_score

Buddy的问题

分配的粒度太大

1页

2页

4页...

我们常常需要分配小内存

能否先分大的，再划分小的等份？




Slab

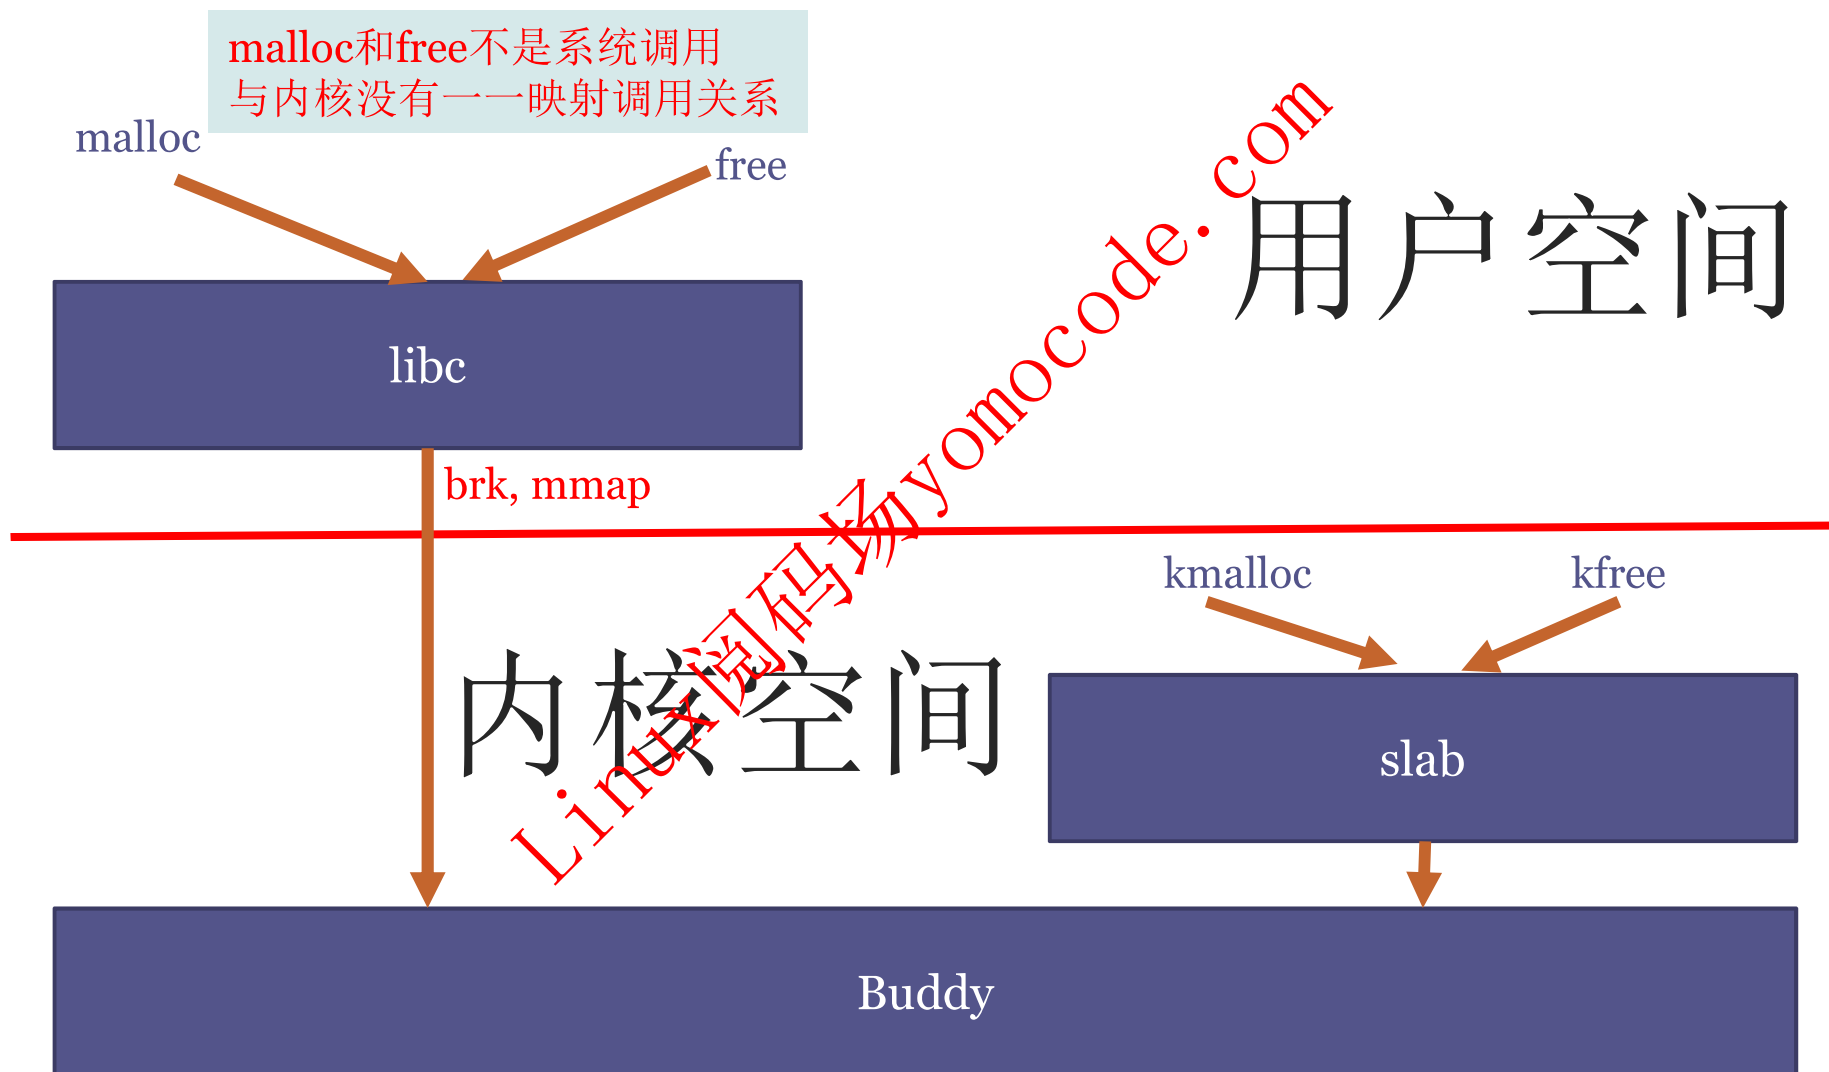
slab

- n. 厚板，平板，厚片；混凝土路面；
- v. 把...分成厚片；用石板铺；

View of a waffle raft
slab being set up,
showing the cardboard
void forms



Libc, Slab 与 buddy



/proc/slabinfo

```
# name          <active_objs> <num_objs> <objsize> <objperslab> <pagesperslab> : tunables <limit> <batchcount>
<active_slabs> <num_slabs> <sharedavail>
isofs_inode_cache 44      44      360      22      2 : tunables 0      0      0 : slabdata 2      2      0
ext4_groupinfo_4k 156     156     104      39      1 : tunables 0      0      0 : slabdata 4      4      0
UDPLITEv6         0        0      768      21      4 : tunables 0      0      0 : slabdata 0      0      0
UDIPv6            84       84      768      21      4 : tunables 0      0      0 : slabdata 4      4      0
tw_sock_TCPv6      0        0      192      21      1 : tunables 0      0      0 : slabdata 0      0      0
TCPv6             44       44     1472     22      8 : tunables 0      0      0 : slabdata 2      2      0
zcache_objnode     0        0      272      30      2 : tunables 0      0      0 : slabdata 0      0      0
kcopyd_job         0        0     2344     13      8 : tunables 0      0      0 : slabdata 0      0      0
dm_uevent          0        0     2464     13      8 : tunables 0      0      0 : slabdata 0      0      0
dm_rq_clone_bio_info 0        0      88      46      1 : tunables 0      0      0 : slabdata 0      0      0
dm_rq_target_io    0        0      264      31      2 : tunables 0      0      0 : slabdata 0      0      0
bsg_cmd            0        0      288      28      2 : tunables 0      0      0 : slabdata 0      0      0
mqueue_inode_cache 28       28      576      28      4 : tunables 0      0      0 : slabdata 1      1      0
fuse_request       42       42      376      21      2 : tunables 0      0      0 : slabdata 2      2      0
fuse_inode         36       36      448      18      2 : tunables 0      0      0 : slabdata 2      2      0
ecryptfs_inode_cache 0        0      640      25      4 : tunables 0      0      0 : slabdata 0      0      0
fat_inode_cache    0        0      416      19      2 : tunables 0      0      0 : slabdata 0      0      0
```

slabtop

```
Active / Total Objects (% used) : 227125 / 230281 (98.6%)
Active / Total Slabs (% used)   : 5040 / 5040 (100.0%)
Active / Total Caches (% used)  : 69 / 100 (69.0%)
Active / Total Size (% used)    : 37463.56K / 38318.42K (97.8%)
Minimum / Average / Maximum Object : 0.01K / 0.17K / 8.00K
```

OBJS	ACTIVE	USE	OBJ SIZE	SLABS	OBJ/SLAB	CACHE SIZE	NAME
37728	37443	99%	0.12K	1179	32	4716K	dentry
29638	29638	100%	0.05K	406	73	1624K	buffer_head
24732	24732	100%	0.59K	916	27	14656K	ext4_inode_cache
22592	22592	100%	0.06K	353	64	1412K	kmalloc-64
17408	17408	100%	0.01K	34	512	136K	ext4_io_page
16128	15529	96%	0.03K	126	128	504K	kmalloc-32
15204	14742	96%	0.09K	362	42	1448K	kmalloc-96
12800	12800	100%	0.01K	25	512	100K	kmalloc-8
8576	8576	100%	0.03K	67	128	268K	anon_vma
8040	8040	100%	0.33K	335	24	2680K	inode_cache
6846	6208	90%	0.19K	326	21	1304K	kmalloc-192
5120	5120	100%	0.02K	20	256	80K	kmalloc-16
3536	2739	77%	0.30K	136	26	1088K	radix_tree_node
1870	1870	100%	0.05K	22	85	88K	Acpi-State
1856	1856	100%	0.06K	29	64	116K	journal_head
1584	1584	100%	0.36K	72	22	576K	proc_inode_cache
1600	1462	91%	0.12K	50	32	200K	kmalloc-128
1328	1245	93%	0.50K	83	16	664K	kmalloc-512
1020	1020	100%	0.02K	6	170	24K	nsproxy
1012	1012	100%	0.36K	46	22	368K	shmem_inode_cache

mallopt

```
#include <malloc.h>
#include <sys/mman.h>

#define SOMESIZE (100*1024*1024) // 100MB
```

```
int main(int argc, char *argv[])
{
```

```
    unsigned char *buffer;
    int i;
```

```
    if (!mlockall(MCL_CURRENT | MCL_FUTURE))
```

```
        mallopt(M_TRIM_THRESHOLD, -1UL);
```

```
    mallopt(M_MMAP_MAX, 0);
```

```
    buffer = malloc(SOMESIZE);
```

```
    if (!buffer)
```

```
        exit(-1);
```

```
    /*
```

```
     * Touch each page in this piece of memory to get it
```

```
     * mapped into RAM
```

```
     */
```

```
    for (i = 0; i < SOMESIZE; i += 4 * 1024)
```

```
        buffer[i] = 0;
```

```
    free(buffer);
```

```
    /* <do your RT-thing> */
```

```
    while(1);
```

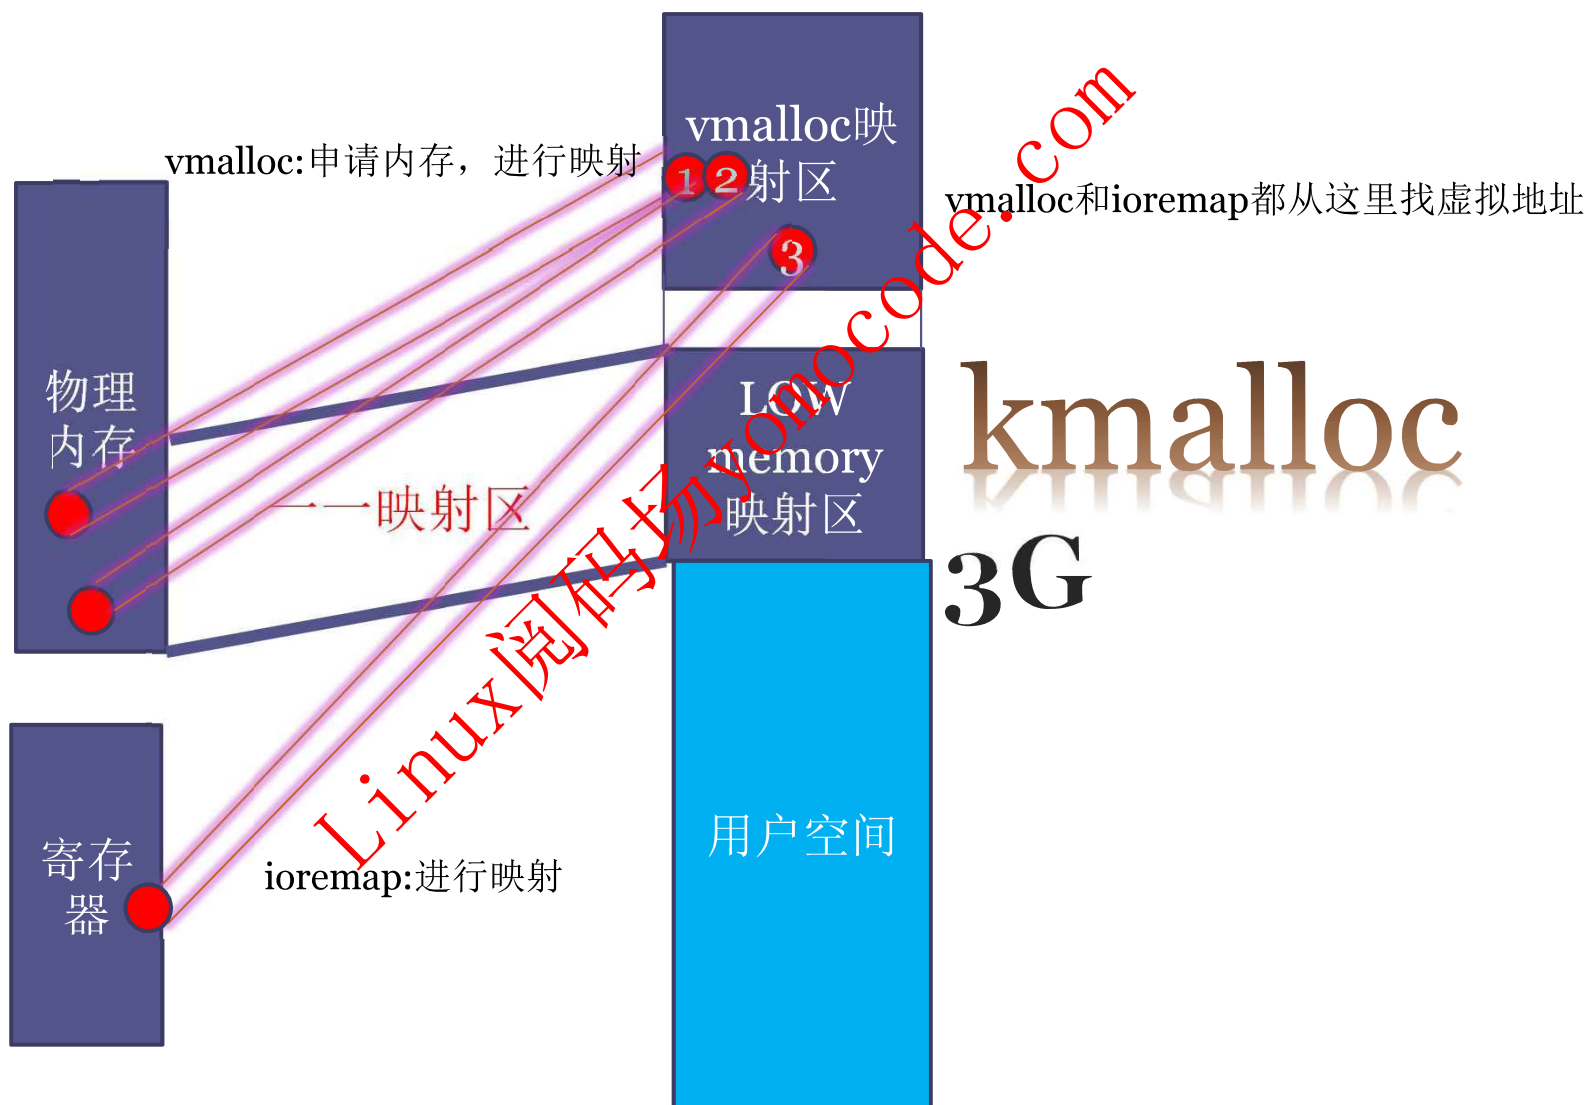
```
    return 0;
```

```
}
```

设置收缩阈值

此时内存不还给内核

kmalloc vs. vmalloc/ioremap




vmallocinfo

从这里可以看出寄存器映射的情况

```
baohua@baohua-VirtualBox:~$ sudo cat /proc/vmallocinfo | grep ioremap
[sudo] password for baohua:
0xf8400000-0xf8404000 16384 acpi_os_map_iomem+0xc7/0x12c phys=3fff0000 ioremap
0xf8404000-0xf8407000 12288 zs_cpu_notifier+0x42/0x80 ioremap
0xf8408000-0xf840b000 12288 zs_cpu_notifier+0x42/0x80 ioremap
0xf840c000-0xf840f000 12288 zs_cpu_notifier+0x42/0x80 ioremap
0xf8410000-0xf8413000 12288 zs_cpu_notifier+0x42/0x80 ioremap
0xf8414000-0xf8417000 12288 pci_iomap_range+0x97/0xe0 phys=f0850000 ioremap
0xf8480000-0xf85b1000 1249280 vesafb_probe+0x482/0x820 phys=e0000000 ioremap
0xf85e0000-0xf85f1000 69632 usb_hcd_pci_probe+0x240/0x620 phys=f0840000 ioremap
0xf8600000-0xf8621000 135168 pci_ioremap_bar+0x38/0x70 phys=f0000000 ioremap
0xf86c0000-0xf86e1000 135168 pci_ioremap_bar+0x38/0x70 phys=f0820000 ioremap
0xf8900000-0xf8d01000 4198400 vboxguestLinuxProbePci+0xbb/0x2f0 [vboxguest] phys=f0400000 ioremap
```

malloc: VSS vs. RSS

1. Program calls `brk()` to grow its heap




2. `brk()` enlarges heap VMA.


New pages are **not** mapped onto physical memory.



3. Program tries to access new memory.
Processor page faults.



4. Kernel assigns page frame to process,
creates PTE, resumes execution. Program is
unaware anything happened.



内存耗尽的程序

```
#include <stdlib.h>
#include <stdio.h>
#include <string.h>
```

```
int main(int argc, char **argv)
{
```

```
    int max = -1;
    int mb = 0;
    char *buffer;
    int i;
```

```
#define SIZE 2000
```

```
    unsigned int *p = malloc(1024 * 1024 * SIZE);
```

```
    printf("malloc buffer: %p\n", p);
```

```
    for (i = 0; i < 1024 * 1024 * (SIZE/sizeof(int)); i++) {
```

```
        p[i] = 123;
```

```
        if ((i & 0xFFFF) == 0) {
```

```
            printf("%dMB written\n", i >> 18);
```

```
            usleep(100000);
```

```
        }
```

```
    }
```

```
    pause();
```

```
    return 0;
```

```
}
```

设定条件

✓ 总内存1G

✓ swapoff -a

✓ echo 1 > /proc/sys/vm/overcommit_memory

写到后面，会OOM

OOM 打 分 因 子

■ mm/oom_kill.c中的badness() 给每个进程一个oom score，取决于：

✓ 驻留内存、 pagetable和swap的使用

采用百分比乘以10（percent-times-ten）：一个使用全部内存的进程得分1000，使用0个字节的进程得分0。

✓ Root用户进程减去30

✓ oom_score_adj: oom_score会加上oom_score_adj这个值

✓ oom_adj: -17~15的系数调整

badness()

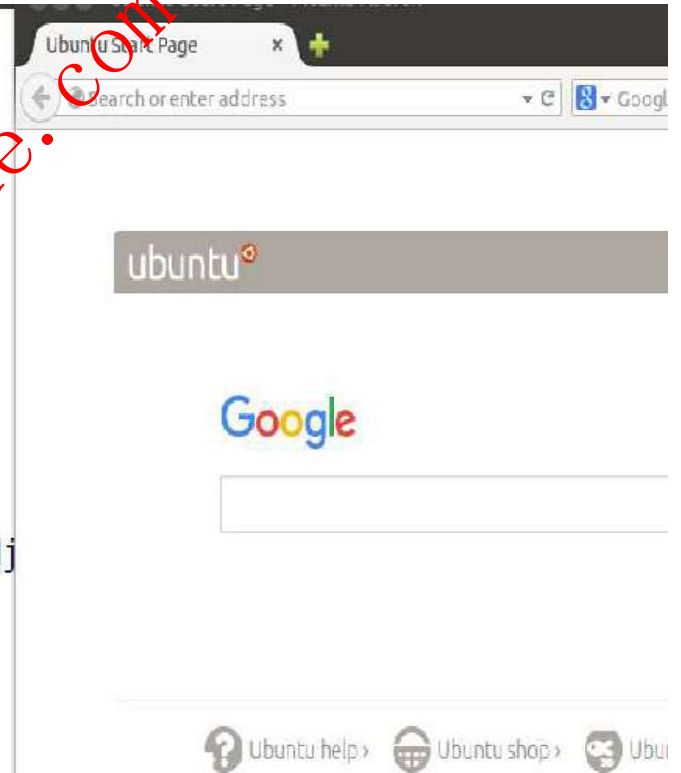
```
/*
 * The baseline for the badness score is the proportion of RAM that each
 * task's rss, pagetable and swap space use.
 */
points = get_mm_rss(p->mm) + get_mm_counter(p->mm, MM_SWAPENTS) +
        atomic_long_read(&p->mm->nr_ptes) + mm_nr_pmds(p->mm);
task_unlock(p);

/*
 * Root processes get 3% bonus, just like the __vm_enough_memory()
 * implementation used by LSMs.
 */
if (has_capability_noaudit(p, CAP_SYS_ADMIN))
    points -= (points * 3) / 100;

/* Normalize to oom_score_adj units */
adj *= totalpages / 1000;
points += adj;
```

通过 oom_adj 调整进程的 oom 打分

```
baohua@baohua-VirtualBox:~$ pidof firefox
3912
baohua@baohua-VirtualBox:~$ cd /proc/3912/
baohua@baohua-VirtualBox:/proc/3912$ echo 1 > oom_adj
baohua@baohua-VirtualBox:/proc/3912$ cat oom_score
144
baohua@baohua-VirtualBox:/proc/3912$ echo 3 > oom_adj
baohua@baohua-VirtualBox:/proc/3912$ cat oom_score
262
baohua@baohua-VirtualBox:/proc/3912$ echo 15 > oom_adj
baohua@baohua-VirtualBox:/proc/3912$ cat oom_score
1084
baohua@baohua-VirtualBox:/proc/3912$ sudo sh -c 'echo -5 > oom_adj'
[sudo] password for baohua:
baohua@baohua-VirtualBox:/proc/3912$ cat oom_score
0
baohua@baohua-VirtualBox:/proc/3912$
```



Android在程序退出时候，并不杀死进程，而是等**OOM**再杀死
Android根据不同的进程类型设置不同的**oom_adj**

课程练习源码

<https://github.com/21cnbao/memory-courses>

Linux阅码场yomocode.com



谢谢！

Linux阅码场yomocode.com