

Linux 任务二 内存管理(三)

Linux 阅码场 yomocode.com

麦当劳喜欢您来，喜欢您再来



扫描关注
Linux阅码场



进程的内存消耗和泄漏

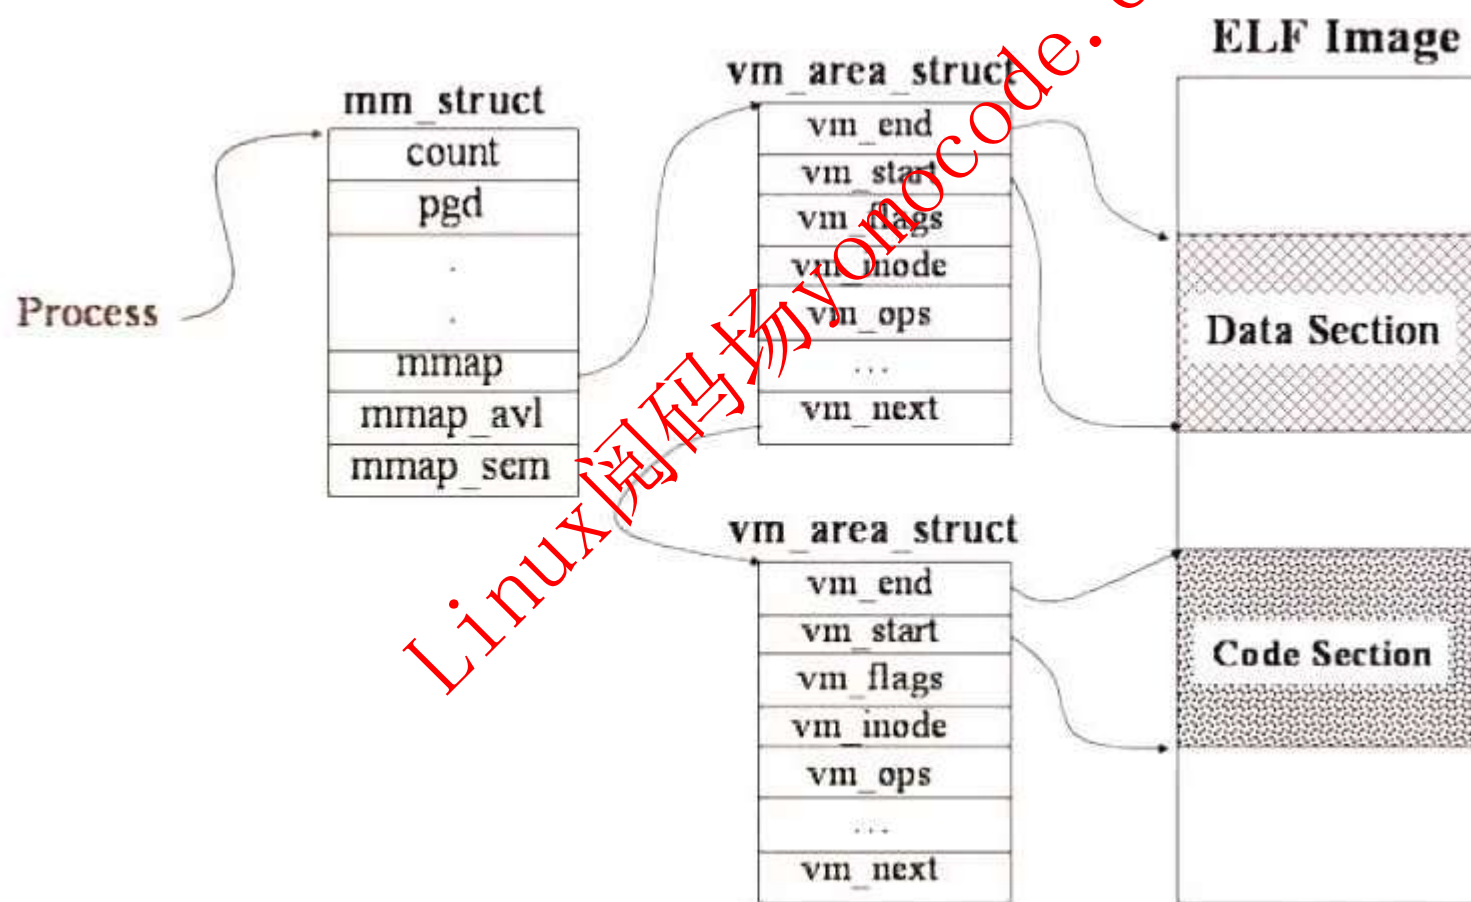
- *进程的VMA。
- *进程内存消耗的4个概念：vss、rss、pss和uss
- *page fault的几种可能性，major和minor。
- *应用内存泄漏的界定方法
- *应用内存泄漏的检测方法：valgrind和addresssanitizer

练习题

- *看一下进程的/proc/<pid>/maps和smaps文件；
- *pmap一个进程；
- *把同一个程序运行2次，运行1次，观察pss；再运行，得到2个进程，观察代码段的pss变化；
- *valgrind检查内存错误

进程的虚拟地址空间VMA

进程的每一段虚拟地址空间就是一个VMA




pmap, /proc/<pid>/maps, /proc/<pid>/smaps

```
baohua@baohua-VirtualBox:~$ pmap 3474
3474:  ./a.out
08048000      4K r-x-- a.out
08049000      4K r---- a.out
0804a000      4K rw--- a.out
b75e0000      4K rw--- [ anon ]
b75e1000    1704K r-x-- libc-2.19.so
b778b000       8K r---- libc-2.19.so
b778d000       4K rw--- libc-2.19.so
b778e000     12K rw--- [ anon ]
b77af000     12K rw--- [ anon ]
b77b2000       8K r---- [ anon ]
b77b4000       8K r-x-- [ anon ]
b77b6000    128K r-x-- ld-2.19.so
b77d6000       4K r---- ld-2.19.so
b77d7000       4K rw--- ld-2.19.so
bfcfb000    132K rw--- [ stack ]
total      2040K
```


```
baohua@baohua-VirtualBox:~$ cat /proc/3474/maps
08048000-08049000 r-xp 00000000 08:01 265913      /home/baohua/a.out
08049000-0804a000 r--p 00000000 08:01 265913      /home/baohua/a.out
0804a000-0804b000 rw-p 00001000 08:01 265913      /home/baohua/a.out
b75e0000-b75e1000 rw-p 00000000 00:00 0
b75e1000-b778b000 r-xp 00000000 08:01 560893      /lib/i386-linux-gnu/libc-2.19.so
b778b000-b778d000 r--p 001aa000 08:01 560893      /lib/i386-linux-gnu/libc-2.19.so
b778d000-b778e000 rw-p 001ac000 08:01 560893      /lib/i386-linux-gnu/libc-2.19.so
b778e000-b7791000 rw-p 00000000 00:00 0
b7791000-b779b2000 rw-p 00000000 00:00 0
b779b2000-b77b2000 r--p 00000000 00:00 0      [vvar]
b77b2000-b77b4000 r-xp 00000000 00:00 0      [vdso]
b77b4000-b77b6000 r-xp 00000000 00:00 0
b77b6000-b77d6000 r-xp 00000000 08:01 575684      /lib/i386-linux-gnu/ld-2.19.so
b77d6000-b77d7000 r--p 0001f000 08:01 575684      /lib/i386-linux-gnu/ld-2.19.so
b77d7000-b77d8000 rw-p 00020000 08:01 575684      /lib/i386-linux-gnu/ld-2.19.so
bfcfb000-bfcfd000 rw-p 00000000 00:00 0      [stack]
```

```
baohua@baohua-VirtualBox:~$ cat /proc/3474/smaps | more
08048000-08049000 r-xp 00000000 08:01 265913      /home/baohua/a.out
Size:                4 kB
Rss:                  4 kB
Pss:                  4 kB
Shared_Clean:         0 kB
Shared_Dirty:         0 kB
Private_Clean:        4 kB
Private_Dirty:        0 kB
Referenced:           4 kB
Anonymous:            0 kB
AnonHugePages:        0 kB
Swap:                 0 kB
KernelPageSize:       4 kB
MMUPageSize:          4 kB
Locked:               0 kB
VmFlags: rd ex mr mw me dw
08049000-0804a000 r--p 00000000 08:01 265913      /home/baohua/a.out
Size:                4 kB
Rss:                  4 kB
```


[illegible]

VSS vs. RSS

1. Program calls `brk()` to grow its heap




2. `brk()` enlarges heap VMA.


New pages are **not** mapped onto physical memory.



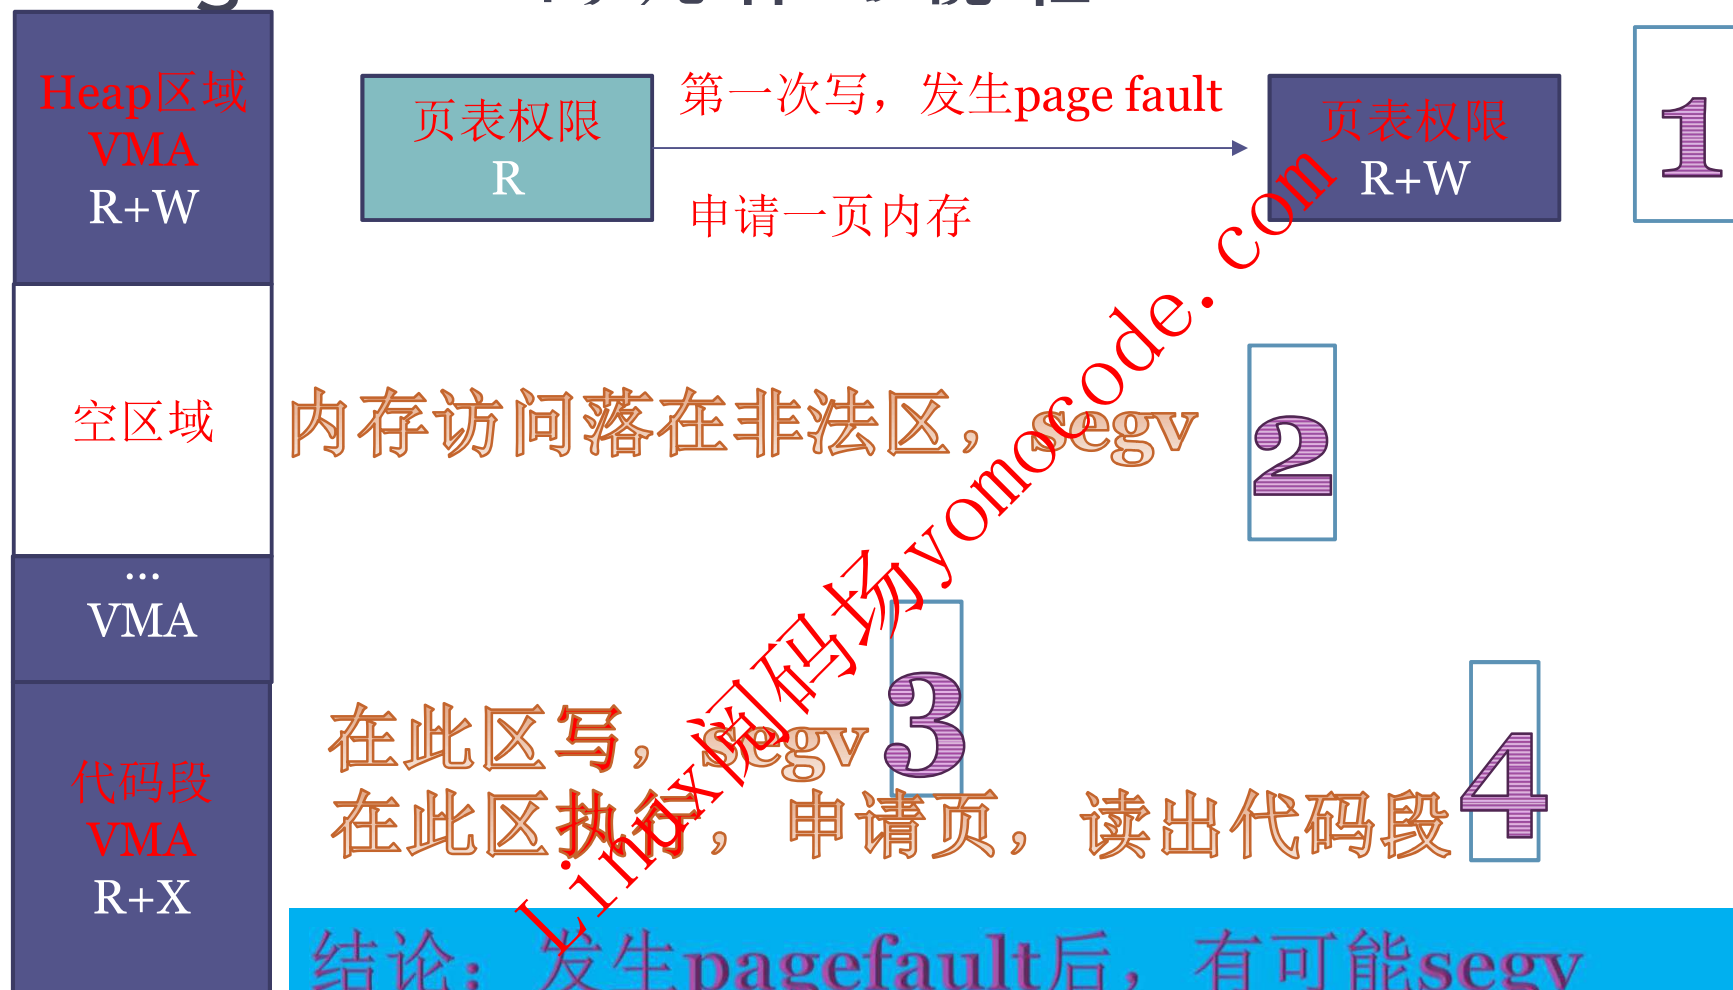
3. Program tries to access new memory.
Processor page faults.



4. Kernel assigns page frame to process,
creates PTE, resumes execution. Program is
unaware anything happened.



Page fault的几种可能性



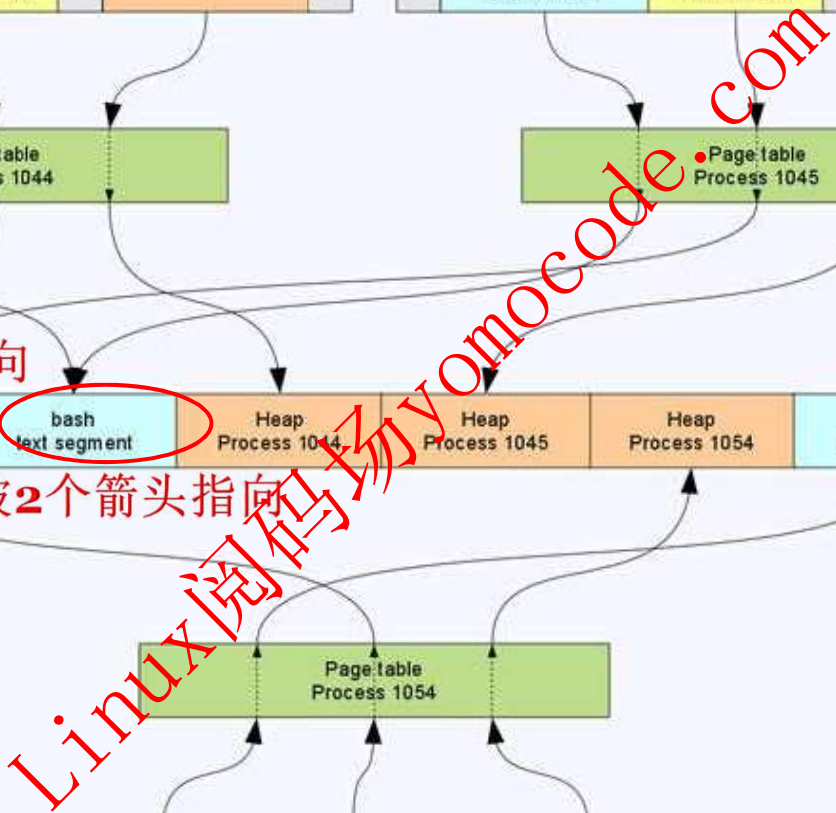
结论: 发生pagefault后, 有可能segv

✓ 非法区域 (2)

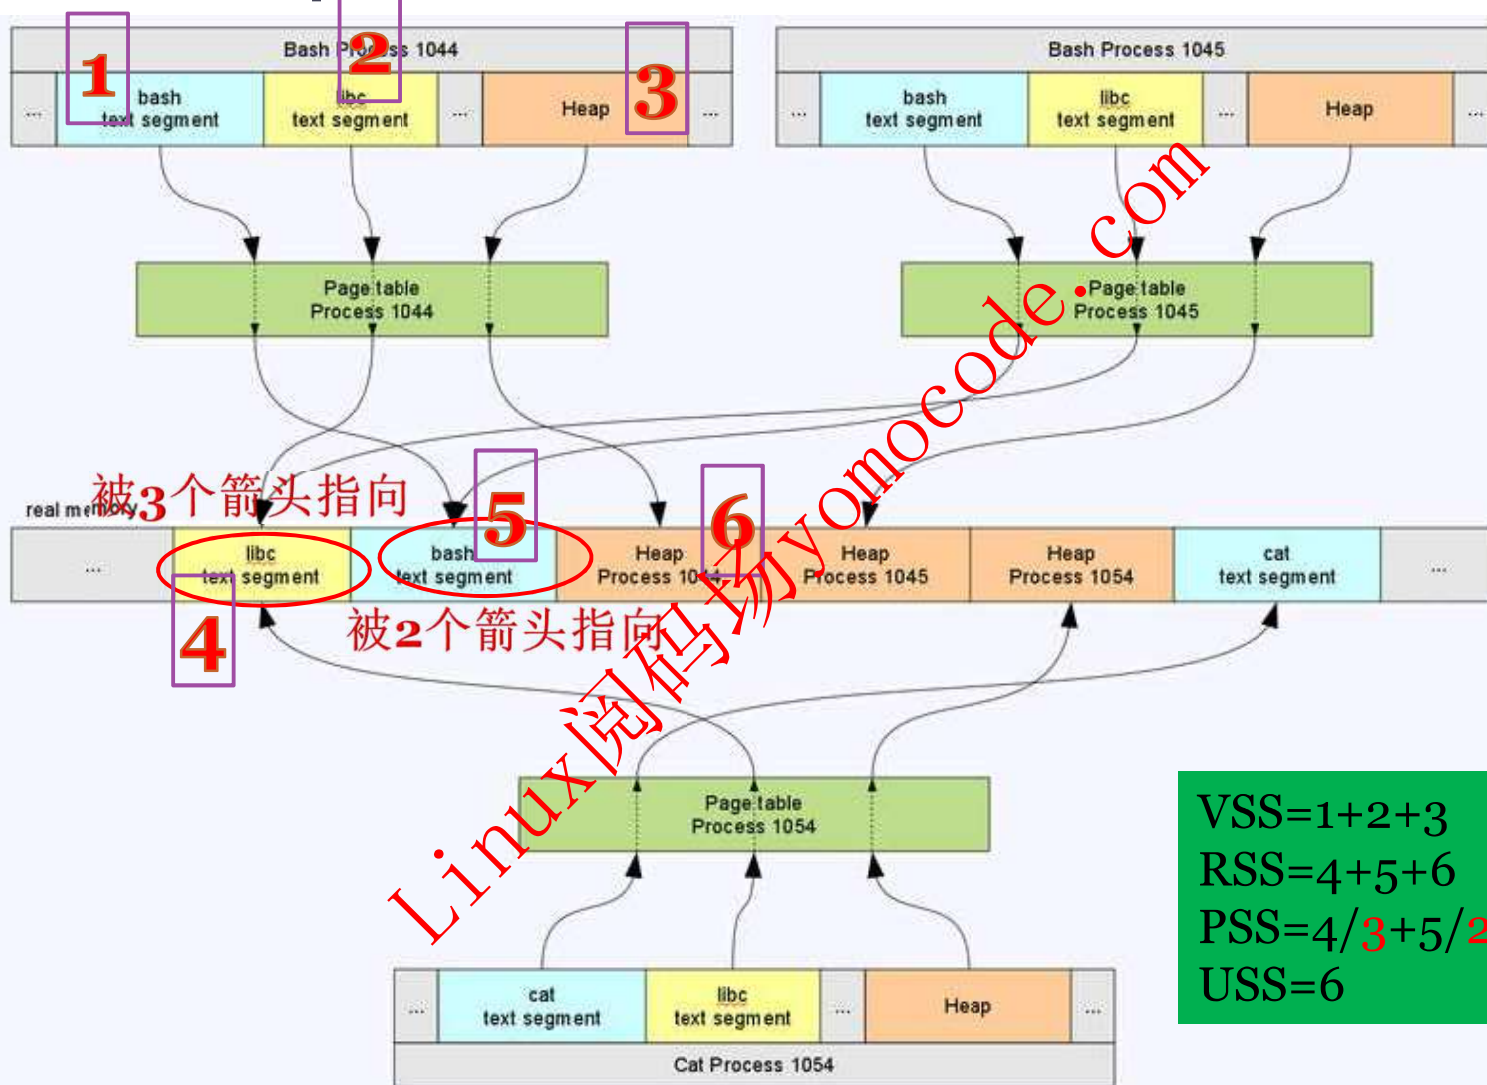
✓ VMA权限不对 (3)

有可能, 不segv, 而是申请内存(1Minor,4Major)

由左被进程如何瓜分？



vss/rss/pss/uss

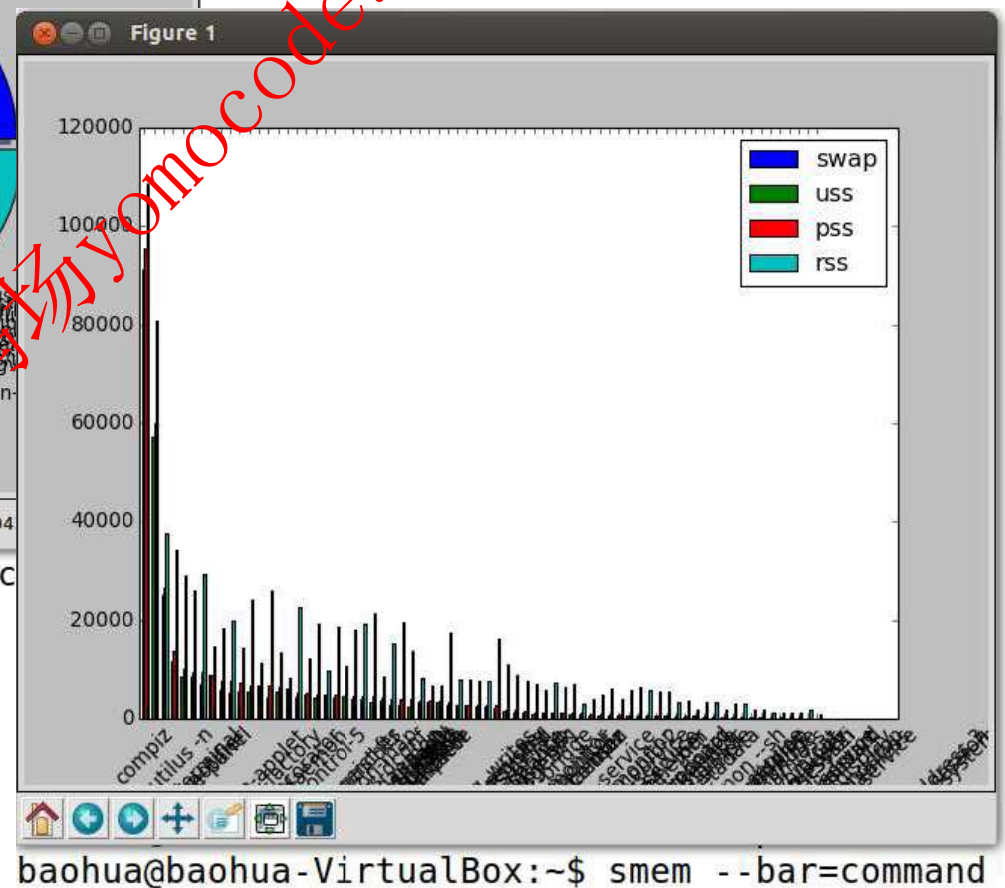



smem – process memory

```
baohua@baohua-VirtualBox:~$ smem
```

PID	User	Command	Swap	USS	PSS	RSS
3474	baohua	./a.out	0	68	76	1028
3352	baohua	/bin/cat	0	112	143	1896
2750	baohua	/usr/bin/VBoxClient --seaml	0	84	217	1164
2743	baohua	/usr/bin/VBoxClient --displ	0	84	219	1172
2755	baohua	/usr/bin/VBoxClient --draga	0	92	219	1140
2733	baohua	/usr/bin/VBoxClient --cliph	0	120	236	1160
2838	baohua	upstart-dbus-bridge --daemo	0	208	271	1696
2839	baohua	upstart-dbus-bridge --daemo	0	228	295	1784
2787	baohua	upstart-event-bridge	0	264	315	2928
3049	baohua	/sbin/initctl emit indicato	0	232	327	2988
2836	baohua	upstart-file-bridge --daemo	0	268	346	1772
2752	baohua	/usr/bin/VBoxClient --seaml	0	244	467	3296
2746	baohua	/usr/bin/VBoxClient --displ	0	248	474	3384
2830	baohua	gpg-agent --daemon --sh	32	528	543	1704
2758	baohua	/usr/bin/VBoxClient --draga	32	332	563	3572
2887	baohua	/bin/dbus-daemon --config-f	0	404	581	3320
3054	baohua	/usr/lib/i386-linux-gnu/ind	0	524	602	5484
3287	baohua	/usr/lib/gvfs/gvfsd-metadat	0	544	618	5580
2892	baohua	/usr/lib/at-spi2-core/at-sp	0	540	640	5844
2960	baohua	/usr/lib/ibus/ibus-engine-s	0	556	691	6364

smem – pie和bar



进程内存泄漏的界定

连续多点采样法，随着时间越久，进程耗费内存越多



观察一个有内存泄漏的进程

```
void main(void)
{
    unsigned int *p1, *p2;
    while(1)
    {
        p1=malloc(4096*3);
        p1[0] = 0;
        p1[1024] = 1;
        p1[1024*2] = 2;

        p2=malloc(1024);
        p2[0] = 1;
        free(p2);
        sleep(1);
    }
}
```

内存随着时间变大

baohua@baohua-VirtualBox:~\$ smem -P a.out

PID	User	Command	Swap	USS	PSS	RSS
3836	baohua	./a.out	0	116	125	1172
3837	baohua	/usr/bin/python /usr/bin/sm	0	4776	4803	6912

baohua@baohua-VirtualBox:~\$ smem -P a.out

PID	User	Command	Swap	USS	PSS	RSS
3836	baohua	./a.out	0	128	137	1184
3838	baohua	/usr/bin/python /usr/bin/sm	0	4776	4801	6840

baohua@baohua-VirtualBox:~\$ smem -P a.out

PID	User	Command	Swap	USS	PSS	RSS
3836	baohua	./a.out	0	140	149	1196
3839	baohua	/usr/bin/python /usr/bin/sm	0	4780	4811	6904

baohua@baohua-VirtualBox:~\$ smem -P a.out

PID	User	Command	Swap	USS	PSS	RSS
3836	baohua	./a.out	0	152	161	1208
3840	baohua	/usr/bin/python /usr/bin/sm	0	4776	4803	6912

内存泄漏的检查-valgrind

valgrind --tool=memcheck --leak-check=yes ./a.out

```
==3978== Memcheck, a memory error detector
==3978== Copyright (C) 2002-2013, and GNU GPL'd, by Julian Seward et al.
==3978== Using Valgrind-3.10.0.SVN and LibVEX; rerun with -h for copyright info
==3978== Command: ./a.out
==3978==
^C==3978==
==3978== HEAP SUMMARY:
==3978==   in use at exit: 73,728 bytes in 6 blocks
==3978==   total heap usage: 12 allocs, 6 frees, 79,872 bytes allocated
==3978==
==3978== 61,440 bytes in 5 blocks are definitely lost in loss record 2 of 2
==3978==    at 0x402A17C: malloc (in /usr/lib/valgrind/vgpreload_memcheck-x86-linux.so)
==3978==    by 0x8048491: main (leak-example.c:6)
==3978==
==3978== LEAK SUMMARY:
==3978==    definitely lost: 61,440 bytes in 5 blocks
==3978==    indirectly lost: 0 bytes in 0 blocks
==3978==    possibly lost: 0 bytes in 0 blocks
==3978==    still reachable: 12,288 bytes in 1 blocks
==3978==    suppressed: 0 bytes in 0 blocks
==3978== Reachable blocks (those to which a pointer was found) are not shown.
```

内存泄漏的检查-addresssanitizer

```
gcc -g -fsanitize=address ./leak-example.c
```

```
./a.out
```

```
=====  
==5232==ERROR: LeakSanitizer: detected memory leaks
```

```
Direct leak of 12288 byte(s) in 1 object(s) allocated from:
```

```
#0 0x7fab49936602 in malloc (/usr/lib/x86_64-linux-gnu/libasan.so.2+0x98602)
```

```
#1 0x4007d7 in main leak-example.c:6
```

```
#2 0x7fab494f482f in __libc_start_main (/lib/x86_64-linux-gnu/libc.so.6+0x2082f)
```

```
SUMMARY: AddressSanitizer: 12288 byte(s) leaked in 1 allocation(s).
```

Linux阅码场yomocode.com

课程练习源码

<https://github.com/21cnbao/memory-courses>

Linux阅码场yomocode.com



谢谢！

Linux阅码场yomocode.com