



[www.yomocode.com](http://www.yomocode.com)

# Linux进程、线程和调度(5)

阅码场™  
[www.yomocode.com](http://www.yomocode.com)

麦当劳喜欢您来，喜欢您再来



扫描关注  
阅码场



# 第五次课大纲

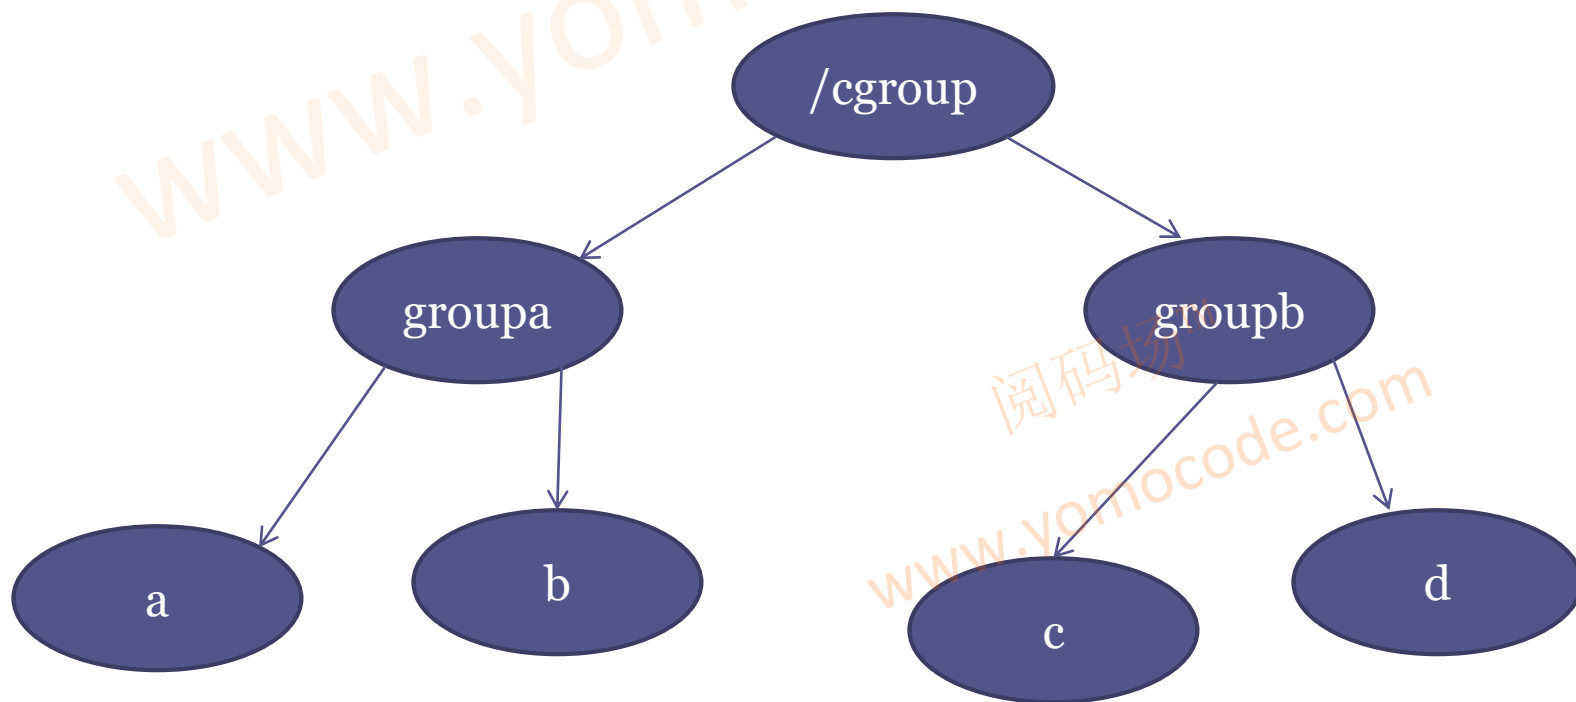
- ❖ 基于cpu cgroups进行CPU资源分配
- ❖ Linux的sched\_autogroup（全新内容）
- ❖ 基于cpuset cgroups进行进程CPU绑定（全新内容）
- ❖ Docker和cgroups
- ❖ Systemd和cgroups（全新内容）
- ❖ Android对cgroups的利用

## 练习题

- ❖ 开启和关闭sched\_autogroup，观察CPU利用率情况；
- ❖ 创建和分群CPU的cgroups，调整权重和quota；
- ❖ 利用cpuset cgroups，进行进程的CPU绑定

# cgroup

- 定义不同cgroup CPU分享的share
- 定义某个cgroup在某个周期里面最多跑多久



# Cgroups 之前的 权重

```
/sys/fs/cgroup/cpu/A$ cat cpu.shares  
1024
```

```
/sys/fs/cgroup/cpu/B$ cat cpu.shares  
2048
```

# Cgroups的资源控制

```
/sys/fs/cgroup/cpu/A
```

```
cpu.cfs_period_us  
cpu.cfs_quota_us
```

阅码场™  
www.yomocode.com

# sched\_autogroup: 200 行的 wonder patch

Yeah. And I have to say that I'm (very happily) surprised by just how small that patch really ends up being, and how it's not intrusive or ugly either.

I'm also very happy with just what it does to interactive performance. Admittedly, my "testcase" is really trivial (reading email in a web-browser, scrolling around a bit, while doing a "make -j64" on the kernel at the same time), but it's a test-case that is very relevant for me. And it is a `_huge_` improvement.

It's an improvement for things like smooth scrolling around, but what I found more interesting was how it seems to really make web pages load a lot faster. Maybe it shouldn't have been surprising, but I always associated that with network performance. But there's clearly enough of a CPU load when loading a new web page that if you have a load average of 50+ at the same time, you `_will_` be starved for CPU in the loading process, and probably won't get all the http requests out quickly enough.

So I think this is firmly one of those "real improvement" patches. Good job. Group scheduling goes from "useful for some specific server loads" to "that's a killer feature".

Linus



# sched\_autogroup: per-session cgroups

```
$ ps -C a.out o pid,ppid,pgid,sid,comm
```

```
PID PPID PGID SID COMMAND
```

```
3832 3778 3832 3778 a.out
```

```
3852 3835 3852 3835 a.out
```

```
$ cat /proc/3832/autogroup
```

```
/autogroup-485 nice 0
```

```
$ cat /proc/3852/autogroup
```

```
/autogroup-487 nice 0
```

```
$ cat /proc/3778/autogroup
```

```
/autogroup-485 nice 0
```

```
$ cat /proc/3835/autogroup
```

```
/autogroup-487 nice 0
```

# Android 和 cgroup

- apps, bg\_non\_interactive

## *Shares:*

apps: cpu.shares = 1024

bg\_non\_interactive: cpu.shares = 52

## *Quota:*

apps:

cpu.rt\_period\_us: 1000000 cpu.rt\_runtime\_us: 800000

bg\_non\_interactive:

cpu.rt\_period\_us: 1000000 cpu.rt\_runtime\_us: 700000

# Docker 和 cgroup

## ■ Docker使用cgroup调配容器的CPU资源

```
$ docker run --cpu-quota 25000 --cpu-period 10000 --cpu-shares 30  
linuxep/lepvo.1
```

```
baohua@ubuntu:~$ docker ps
```

CONTAINER ID	IMAGE	COMMAND	CREATED
STATUS	PORTS	NAMES	
3f39ca25d14d			

```
baohua@ubuntu:/sys/fs/cgroup/cpu/docker$ cd 3f39c...
```

```
baohua@ubuntu:/sys/fs/cgroup/cpu/docker/3f39c...$ ls
```

```
cgroup.clone_children cgroup.procs cpuacct.stat cpuacct.usage  
cpuacct.usage_percpu cpu.cfs_period_us cpu.cfs_quota_us cpu.shares cpu.stat  
notify_on_release tasks
```

```
baohua@ubuntu:/sys/fs/cgroup/cpu/docker/3f39c...$ cat cpu.cfs_quota_us
```

```
25000
```

```
baohua@ubuntu:/sys/fs/cgroup/cpu/docker/3f39c...$ cat cpu.cfs_period_us
```

```
10000
```

```
baohua@ubuntu:/sys/fs/cgroup/cpu/docker/3f39c...$ cat cpu.shares
```

```
30
```

# systemd cg 层级

- ❖ slice
- ❖ scope
- ❖ service

Control group /:

```
-.slice
├── user.slice
│   ├── user-1000.slice
│   │   ├── user@1000.service
│   │   │   ├── gvfs-goa-volume-monitor.service
│   │   │   │   └── 1944 /usr/lib/gvfs/gvfs-goa-volume-monitor
│   │   │   ├── init.scope
│   │   │   │   ├── 1502 /lib/systemd/systemd --user
│   │   │   │   └── 1521 (sd-pam)
│   │   │   ├── gvfs-gphoto2-volume-monitor.service
│   │   │   │   └── 1940 /usr/lib/gvfs/gvfs-gphoto2-volume-monitor
│   │   │   ├── at-spi-dbus-bus.service
│   │   │   │   ├── 1763 /usr/lib/at-spi2-core/at-spi-bus-launcher
│   │   │   │   ├── 1768 /usr/bin/dbus-daemon --config-file=/usr/share/defaults/at-spi2...
│   │   │   │   └── 1770 /usr/lib/at-spi2-core/at-spi2-registryd --use-gnome-session
│   │   └── init.scope
│   │       └── 1 /sbin/init splash
│   └── system.slice
│       ├── irqbalance.service
│       │   └── 621 /usr/sbin/irqbalance --foreground
│       ├── ssh.service
│       │   └── 791 /usr/sbin/sshd -D
│       ├── systemd-logind.service
│       │   └── 605 /lib/systemd/systemd-logind
```

# systemd cg 相关命令

- systemd-cgls
- systemd-cgtop
- systemd-run
- systemctl

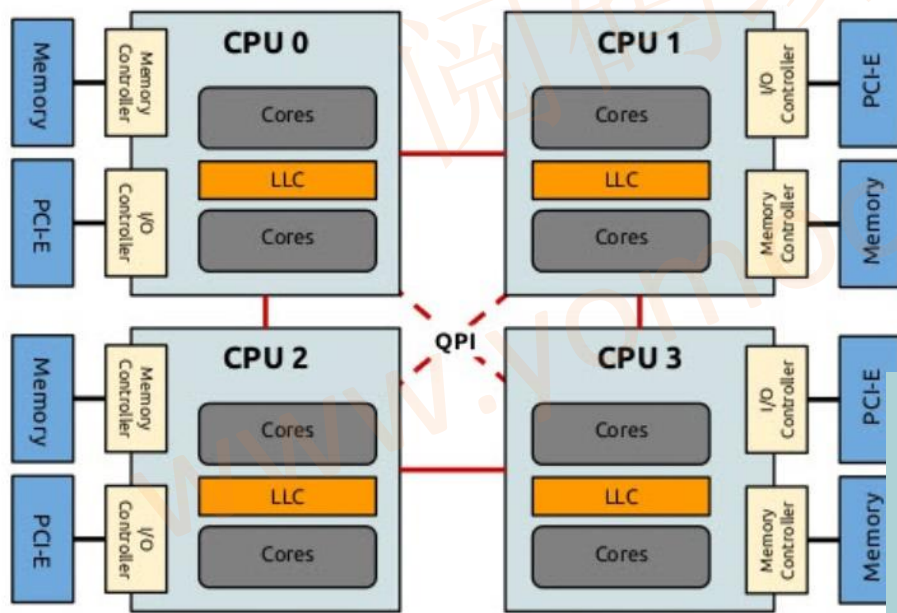
set-property 限制cgroup的资源

# cpuset

- Cpusets provide a Linux kernel mechanism to constrain which CPUs and Memory Nodes are used by a process or set of processes.
- The root cpuset contains all the systems CPUs and Memory Nodes.
- `cpuset.cpus`: list of CPUs in that cpuset
- `cpuset.mems`: list of Memory Nodes in that cpuset

# NUMA

CPU architecture (Intel Sandy Bridge)



```
# numactl --hardware
available: 2 nodes (0-1)
node 0 cpus: 0 1 2 3 4 5 12 13 14 15 16 17
node 0 size: 130677 MB
node 0 free: 1453 MB
node 1 cpus: 6 7 8 9 10 11 18 19 20 21 22 23
node 1 size: 131056 MB
node 1 free: 614 MB
node distances:
node  0  1
   0: 10 21
   1: 21 10
```

谢谢!

阅码场™

www.yomocode.com

阅码场™

www.yomocode.com





阅码场出品

[www.yomocode.com](http://www.yomocode.com)